

16-11-2020

CosmoML 2020/21 - 3rd meeting

- Recap
- Evaluation metrics
- Decision Trees, Random Forests, Support Vector Machines, k Nearest Neighbors
- Hyperparameter tuning+pipelines in sklearn



Michele Mancarella

Recap & classification metrics



Supervised ML workflow

Data

$$\mathcal{D} = \{(\mathbf{X}, y)_i\}_{i=1}^N$$


Model $f(\mathbf{X}; \mathbf{w})$ + cost $\mathcal{C}(y, f(\mathbf{X}; \mathbf{w}))$

Must be differentiable wrt weights!


Train

$$\hat{\mathbf{w}} = \arg \min_{\mathbf{w}} \mathcal{C}(y, f(\mathbf{X}; \mathbf{w})) \quad \text{by gradient descent}$$

Split dataset



Hyperparameter tuning

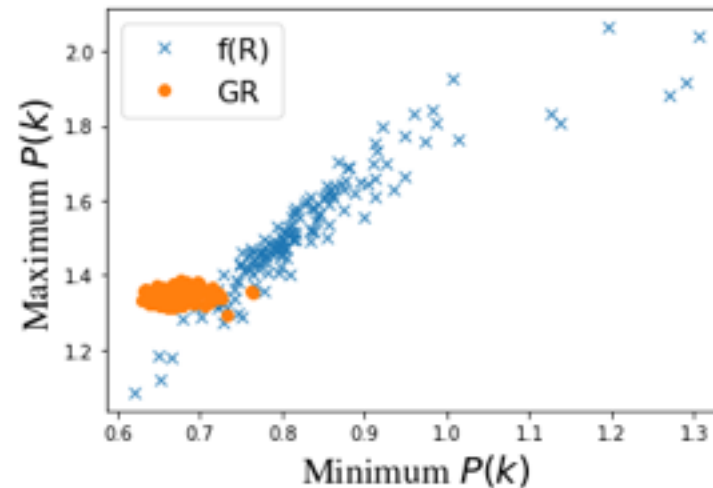


Test - Evaluation metric

In general evaluation metric \neq cost,
not differentiable wrt weights

Classification problems

Data



$$\mathcal{D} = \{(\mathbf{X}, y)_i\}_{i=1}^N$$

$$y = fR, GR$$

$$\mathbf{X} = (\min P, \max P) (+ \text{ powers...})$$

Model $f(\mathbf{X}; \mathbf{w})$

$$\sigma(\mathbf{w}^T \mathbf{X}), \quad \sigma(t) = \frac{1}{1 + e^{-t}}$$

Cost $\mathcal{C}(y, f(\mathbf{X}; \mathbf{w}))$

binomial likelihood/ binary cross-entropy loss.

$$\mathcal{C} = -\log \mathcal{L}(\mathcal{D}; \mathbf{w}) = -\sum_i y_i \log f(\mathbf{w}^T \mathbf{x}_i) - (1 - y_i) \log[1 - f(\mathbf{w}^T \mathbf{x}_i)]$$

Train

Test /validation

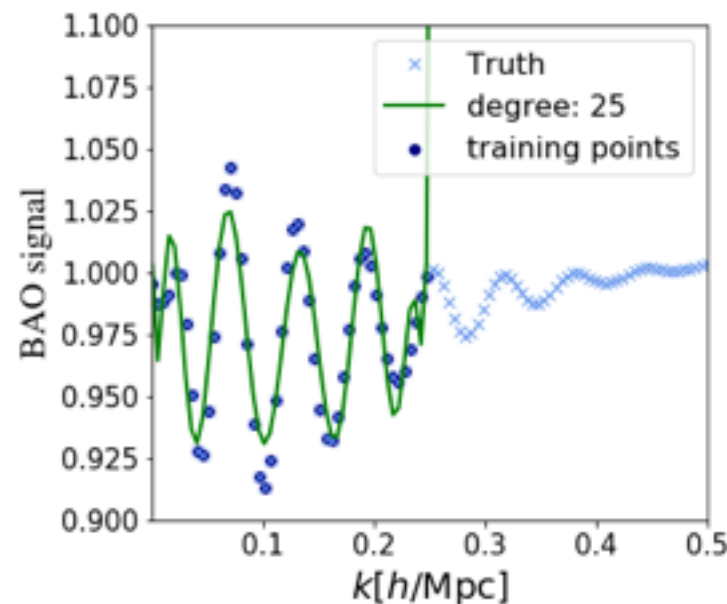
Set threshold prob.

Get sigmoid output & assign label depending on threshold

Evaluation metric = accuracy

Regression problems

Data



y = BAO signal
 $X = k$ (+ powers...)

Model $f(\mathbf{X}; \mathbf{w})$ $f(k; \mathbf{w}) = w_0 + w_1 k + w_2 k^2 + \dots$

Cost $\mathcal{C}(y, f(\mathbf{X}; \mathbf{w}))$

MSE / gaussian likelihood

$$\mathcal{C} = -\log \mathcal{L}(\mathcal{D}; \mathbf{w}) \propto -\sum_i (y - f)^2$$

Train

Test /validation

Evaluation metric = cost

Metrics for classification

Classification problems: Confusion matrix (not a “metric”, but check it...)

	<u>Actual</u>	
	1	0
<u>Predicted</u>	1 True positive	0 False positive
	0 False negative	0 True negative

LCDM=0

fR=1 (or: has covid-19)

ACCURACY= % of correctly predicted.

!! Careful with unbalanced datasets/type of problem. E.g. if I want to be sure to catch an fR in a dataset with 99 LCDM and 1 fR, a classifier that always predicts LCDM would have a 99% accuracy!
(or fraud/spam detection, covid-19 tests...)


RECALL = fraction of samples from a class which are correctly predicted
(e.g. precision on fR: $TP/(TP+FN)$
how many that actually have covid are caught?)

“Well then let’s predict all positive all the time!”

PRECISION = fraction of samples predicted in a class which are actually in that class
(e.g. precision on fR: $TP/(TP+FP)$
how many that are predictive positive are actually positive?)

Metrics


ROC CURVE



ML Algorithms



Decision trees



Tree = Oriented graph w. any two nodes connected by 1 edge

Algorithm 1 Decision Tree

Start at root node

repeat

\forall partition $S(\theta)$, $\theta = (j, t_j)$ consisting of feature j and threshold t_j of data at node k :


Compute *impurity* $I(S, \theta) = \frac{n_{\text{left}}}{N_k} H(S_{\text{left}}) + \frac{n_{\text{right}}}{N_k} H(S_{\text{right}})$

Choose the partition corresponding to $\hat{\theta} = \arg \min_{\theta} I(S, \theta)$

until Max depth is reached or $N_k = N_{\text{min}}$

S_{left} = partition of data w. feature $j < t_j$, $S_{\text{right}} = S \setminus S_{\text{left}}$

Decision trees



Gini

$$H(x_k) = \sum_{i=1}^{N_{classes}} p_{ik}(1 - p_{ik})$$

Entropy

$$H(x_k) = - \sum_{i=1}^{N_{classes}} p_{ik} \log p_{ik}$$

p_{ik} = fraction of points in node k in class i

(For regression problems: MSE)

`sklearn.tree.DecisionTreeClassifier`

```
class sklearn.tree.DecisionTreeClassifier(criterion='gini', splitter='best', max_depth=None, min_samples_split=2, min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features=None, random_state=None, max_leaf_nodes=None, min_impurity_decrease=0.0, min_impurity_split=None, class_weight=None, presort='deprecated', ccp_alpha=0.0) \[source\]
```

A decision tree classifier.


Decision trees

Advantages:

- ♦ No feature preparation - also w. mixed numerical&categorical
- ♦ “**White box**” ! (Visualizable, interpretable, simple). Feature importance:

$$\frac{n_k}{N} \left(I(S) - \frac{n_{\text{left}}}{N_k} I(S_{\text{left}}) - \frac{n_{\text{right}}}{N_k} I(S_{\text{right}}) \right)$$

$$FI_{f_i} = \frac{\sum_{j \text{ s.t. node } j \text{ splits on } i\text{-th feature}} GI_j}{\sum_j GI_j}$$



Disadvantages:

- ♦ Easily overfit
 - ♦ Handling of unbalanced classes
 - ♦ Unstable to variations in training data
 - ♦ Greedy ! (local minima)
- ← REGULARIZATION: max depth, max n. of leafs, min. sample split


Ensemble methods

- ❖ Random Forests - grow different trees w. bootstrap
- ❖ AdaBoost - re-weight error to reinforce points where classifier performs poorly (i.e. concentrate on difficult examples)
- ❖ Gradient Boosted Trees/XGBoost

Random Forests

Add randomness to prevent overfit (everywhere in ML, cfr dropout in NN)

- (1) Grow different trees on bootstrap sample
- (2) When splitting each node, partition using a random subset of features



3.2.4.3.1. `sklearn.ensemble.RandomForestClassifier`


```
class sklearn.ensemble. RandomForestClassifier(n_estimators=100, criterion='gini', max_depth=None,  
min_samples_split=2, min_samples_leaf=1, min_weight_fraction_leaf=0.0, max_features='auto', max_leaf_nodes=None,  
min_impurity_decrease=0.0, min_impurity_split=None, bootstrap=True, oob_score=False, n_jobs=None, random_state=None,  
verbose=0, warm_start=False, class_weight=None, ccp_alpha=0.0, max_samples=None)
```

[\[source\]](#)

A random forest classifier.


Support Vector Machines

Basic idea: find optimal separating hyperplane (then generalize to non-linear boundaries)




Support Vector Machines

Basic idea: find optimal separating hyperplane (then generalize to non-linear boundaries)




Support Vector Machines

Basic idea: find optimal separating hyperplane (then generalize to non-linear boundaries)




Support Vector Machines

Basic idea: find optimal separating hyperplane (then generalize to non-linear boundaries)



Support Vector Machines

Basic idea: find optimal separating hyperplane (then generalize to non-linear boundaries)




Intuitive formulation:

$$\begin{aligned} \max_{\gamma, w, b} \quad & \gamma \\ \text{s.t.} \quad & y^{(i)}(w^T x^{(i)} + b) \geq \gamma, \quad i = 1, \dots, m \\ & \|w\| = 1. \end{aligned}$$

- Tough constraint to solve
- Far away points should not count

Support Vector Machines

Basic idea: find optimal separating hyperplane (then generalize to non-linear boundaries)




Less intuitive but optimal formulation:

$$\begin{aligned} \max_{\alpha} \quad & W(\alpha) = \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y^{(i)} y^{(j)} \alpha_i \alpha_j \langle x^{(i)}, x^{(j)} \rangle. \\ \text{s.t.} \quad & \alpha_i \geq 0, \quad i = 1, \dots, m \\ & \sum_{i=1}^m \alpha_i y^{(i)} = 0, \end{aligned}$$

- Use Lagrange multipliers α
- Solve the **dual** problem, i.e. maximize over α subject to relations implied by the constraints for \mathbf{w} and \mathbf{b} instead of maximizing over \mathbf{w} and \mathbf{b} subject to the constraint involving α

Support Vector Machines

Basic idea: find optimal separating hyperplane (then generalize to non-linear boundaries)



Less intuitive but optimal formulation:

$$\begin{aligned} \max_{\alpha} \quad & W(\alpha) = \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y^{(i)} y^{(j)} \alpha_i \alpha_j \langle x^{(i)}, x^{(j)} \rangle. \\ \text{s.t.} \quad & \alpha_i \geq 0, \quad i = 1, \dots, m \\ & \sum_{i=1}^m \alpha_i y^{(i)} = 0, \end{aligned}$$


Only true for SV !!

scalar product

- Use Lagrange multipliers α
- Solve the **dual** problem, i.e. maximize over α subject to relations implied by the constraints for **w** and **b** instead of maximizing over **w** and **b** subject to the constraint involving α

Support Vector Machines

Basic idea: find optimal separating hyperplane (then generalize to non-linear boundaries)



Less intuitive but optimal formulation:

$$\begin{aligned} \max_{\alpha} \quad & W(\alpha) = \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y^{(i)} y^{(j)} \alpha_i \alpha_j \langle x^{(i)}, x^{(j)} \rangle. \\ \text{s.t.} \quad & \alpha_i \geq 0, \quad i = 1, \dots, m \\ & \sum_{i=1}^m \alpha_i y^{(i)} = 0, \end{aligned}$$

Only true for SV !!

scalar product

Non-linear boundaries

- Replace scalar product w. $\langle \phi(x), \phi(z) \rangle$ for any non-linear mapping
- Or directly define a KERNEL $K(x, z) = \phi(x)^T \phi(z)$

Support Vector Machines

Regularization:

$$\begin{aligned} \max_{\alpha} \quad & W(\alpha) = \sum_{i=1}^m \alpha_i - \frac{1}{2} \sum_{i,j=1}^m y^{(i)} y^{(j)} \alpha_i \alpha_j \langle x^{(i)}, x^{(j)} \rangle \\ \text{s.t.} \quad & 0 \leq \alpha_i \leq C, \quad i = 1, \dots, m \\ & \sum_{i=1}^m \alpha_i y^{(i)} = 0, \end{aligned}$$

Advantages:

- ♦ Nicely generalizes to non-linear boundaries, versatile
- ♦ Good for higher-dimensional problems
- ♦ Quadratic optimization problem

Disadvantages:

- ♦ No probability estimates!
- ♦ Don't forget normalization !

k-Nearest Neighbors

Basic idea: assign as label the most frequent label among the closest k point (k nearest neighbors) in the feature space

- ✦ Not a minimization problem/nonparametric : simply stores instances of training data
- ✦ Can tune k and distance



Algorithm 1 kNN

for $i = 1, \dots, N_{new}$ **do** Compute distance $d(X_i, x_{new})$
end for

Find set S of points with k smallest distances $d(X_i, x_{new})$, $i \in S$

return Majority label in S

Advantages:

- ✦ Simple but effective
- ✦ Can handle complex boundaries

Disadvantages:

- ✦ No probability estimates!
- ✦ Don't forget normalization !
- ✦ $O(N^2)$ if you don't use smart algorithms