

# Machine Learning 4:

## Neural Networks



Consider a face recognition problem




Consider a face recognition problem




For a  $50 \times 50$  grayscale image, we have 2500 basic features:

$$x_{1,\dots,2500} \in \{0, 255\}$$

As before, we can map in feature space the 'face' and 'non-face' labels




As before, we can map in feature space the 'face' and 'non-face' labels



We definitely require polynomial features to map this decision boundary ...

As before, we can map in feature space the 'face' and 'non-face' labels



We definitely require polynomial features to map this decision boundary ...  
Up to quadratic order this is already  $\sim 3,000,000$  features ....





For a  $2000 \times 2000$  colour image, we have  $2000^2$  basic features ...





For a  $2000 \times 2000$  colour image, we have  $2000^2$  basic features ...

... and we probably want at least up to cubic order for a low bias fit to the data ...



For a  $2000 \times 2000$  colour image, we have  $2000^2$  basic features ...

... and we probably want at least up to cubic order for a low bias fit to the data ...




For a  $2000 \times 2000$  colour image, we have  $2000^2$  basic features ...

... and we probably want at least up to cubic order for a low bias fit to the data ...


... you get the picture

# Neural Networks

A simple neural network




# Neural Networks




- A basic neural network can be seen as multi-layered, logistic regression based, map, or a string of **logistic units**.

# Neural Networks



- A basic neural network can be seen as multi-layered, logistic regression based, map, or a string of **logistic units**.
- Neural networks provide a means to create 'non-linear' features from a set of basic input features. These are called **'activations'**.

# Neural Networks



- A basic neural network can be seen as multi-layered, logistic regression based, map, or a string of **logistic units**.
- Neural networks provide a means to create 'non-linear' features from a set of basic input features. These are called **'activations'**.

Let's see what this means more concretely

For any operation with a well defined goal, one can make the reasonable proposition that this goal can be achieved through some number of logical operations. This is arguably the basis of a simple neural network.




For any operation with a well defined goal, one can make the reasonable proposition that this goal can be achieved through some number of logical operations. This is arguably the basis of a simple neural network.

Consider the following:

For any operation with a well defined goal, one can make the reasonable proposition that this goal can be achieved through some number of logical operations. This is arguably the basis of a simple neural network.

Consider the following:




$a_1$	$a_2$	$h = \text{NAND}$
0	0	1
0	1	1
1	0	1
1	1	0

$$\begin{aligned}
 h(\theta \underline{a}) &= \text{sig}(\theta_0 a_0 + \theta_1 a_1 + \theta_2 a_2) \\
 &= \text{sig}(30 - 20a_1 - 20a_2)
 \end{aligned}$$

For any operation with a well defined goal, one can make the reasonable proposition that this goal can be achieved through some number of logical operations. This is arguably the basis of a simple neural network.


Consider the following:




$a_1$	$a_2$	$h$ = NAND
0	0	1
0	1	1
1	0	1
1	1	0

$$\begin{aligned}
 h(\theta \underline{a}) &= \text{sig}(\theta_0 a_0 + \theta_1 a_1 + \theta_2 a_2) \\
 &= \text{sig}(30 - 20a_1 - 20a_2)
 \end{aligned}$$

Let's now consider some binary classification as our goal  
eg. face or not face.




How is the hypothesis defined?



How is the hypothesis defined?

$$h(\Theta, \mathbf{x}) = a_1^3 = \text{sig}((\theta^{(2)})^T \mathbf{a}),$$

$$(\theta^{(2)})^T \mathbf{a} = \theta_{10}^{(2)} a_0^2 + \theta_{11}^{(2)} a_1^2 + \theta_{12}^{(2)} a_2^2 + \theta_{13}^{(2)} a_3^2 \quad (1)$$



How is the hypothesis defined?

$$h(\Theta, \mathbf{x}) = a_1^2 = \text{sig}((\theta^{(2)})^T \mathbf{a}),$$

$$(\theta^{(2)})^T \mathbf{a} = \theta_{10}^{(2)} a_0^1 + \theta_{11}^{(2)} a_1^1 + \theta_{12}^{(2)} a_2^1 + \theta_{13}^{(2)} a_3^1 \quad (1)$$

where

$$\tilde{\mathbf{a}} = [a_1^1, a_2^1, a_3^1] = \text{sig}((\theta^{(1)})^T \mathbf{x}),$$

$$a_i^1 = \text{sig}(\theta_{i0}^{(1)} x_0 + \theta_{i1}^{(1)} x_1 + \theta_{i2}^{(1)} x_2) \quad (2)$$

## Some points

- The bias units  $x_0 = a_0^2 = a_0^3 = \dots = 1$  allow for an intercept  $\theta_{10}$  in each input.

## Some points

- The bias units  $x_0 = a_0^2 = a_0^3 = \dots = 1$  allow for an intercept  $\theta_{10}$  in each input.
- To compute our prediction  $h(\Theta, \mathbf{x})$ , we must compute all the activation units of each layer in a process called **forward propagation**.



## Some points

- The bias units  $x_0 = a_0^2 = a_0^3 = \dots = 1$  allow for an intercept  $\theta_{10}$  in each input.
- To compute our prediction  $h(\Theta, \mathbf{x})$ , we must compute all the activation units of each layer in a process called **forward propagation**.
- The structure of a neural network, i.e. number of layers and activation units per layer, is called the **architecture**.

## Some points

- The bias units  $x_0 = a_0^2 = a_0^3 = \dots = 1$  allow for an intercept  $\theta_{10}$  in each input.
- To compute our prediction  $h(\Theta, \mathbf{x})$ , we must compute all the activation units of each layer in a process called **forward propagation**.
- The structure of a neural network, i.e. number of layers and activation units per layer, is called the **architecture**.
- We can generalise to **multiclass classification** by introducing more units in the output layer (ex.  $a_2^3, a_3^3$  etc ). We then just choose the output unit with the highest probability for our prediction.

## Some points

- The bias units  $x_0 = a_0^2 = a_0^3 = \dots = 1$  allow for an intercept  $\theta_{10}$  in each input.
- To compute our prediction  $h(\Theta, \mathbf{x})$ , we must compute all the activation units of each layer in a process called **forward propagation**.
- The structure of a neural network, i.e. number of layers and activation units per layer, is called the **architecture**.
- We can generalise to **multiclass classification** by introducing more units in the output layer (ex.  $a_2^3, a_3^3$  etc ). We then just choose the output unit with the highest probability for our prediction.
- We want to find  $\Theta$  which is a collection of weight vectors/matrices  $\theta^j$  where 'j' is the number of layers. Note that the dimension of  $\theta^j$  is  $s_{j+1} \times (s_j + 1)$  where  $s_j$  is the number of units in layer  $j$  excluding the bias unit.

## Some points

- The bias units  $x_0 = a_0^2 = a_0^3 = \dots = 1$  allow for an intercept  $\theta_{10}$  in each input.
- To compute our prediction  $h(\Theta, \mathbf{x})$ , we must compute all the activation units of each layer in a process called **forward propagation**.
- The structure of a neural network, i.e. number of layers and activation units per layer, is called the **architecture**.
- We can generalise to **multiclass classification** by introducing more units in the output layer (ex.  $a_2^3, a_3^3$  etc ). We then just choose the output unit with the highest probability for our prediction.
- We want to find  $\Theta$  which is a collection of weight vectors/matrices  $\theta^j$  where 'j' is the number of layers. Note that the dimension of  $\theta^j$  is  $s_{j+1} \times (s_j + 1)$  where  $s_j$  is the number of units in layer  $j$  excluding the bias unit.

So what is our cost function ?

For a neural network with  $\mathbf{L}$  layers, looking to classify into  $\mathbf{K}$  different classes, and a training set of  $\mathbf{m}$  examples, we can extend our old logistic cost function as follows

$$J(\Theta) = -\frac{1}{m} \sum_{i=1}^m \sum_{k=1}^K [y_k^i \log[h_k(x^i)] + (1 - y_k^i) \log[1 - h_k(x^i)]] \\ + \frac{\lambda}{2m} \sum_{\ell=1}^{L-1} \sum_{i=1}^{s_\ell} \sum_{j=1}^{s_{\ell+1}} (\theta_{ji}^\ell)^2, \quad (3)$$

where  $\lambda$  is our regularisation parameter and  $s_\ell$  is the number of units in the  $\ell^{th}$  layer **not** including the bias unit.

For a neural network with  $\mathbf{L}$  layers, looking to classify into  $\mathbf{K}$  different classes, and a training set of  $\mathbf{m}$  examples, we can extend our old logistic cost function as follows

$$J(\Theta) = -\frac{1}{m} \sum_{i=1}^m \sum_{k=1}^K [y_k^i \log[h_k(x^i)] + (1 - y_k^i) \log[1 - h_k(x^i)]] \\ + \frac{\lambda}{2m} \sum_{\ell=1}^{L-1} \sum_{i=1}^{s_\ell} \sum_{j=1}^{s_{\ell+1}} (\theta_{ji}^\ell)^2, \quad (3)$$

where  $\lambda$  is our regularisation parameter and  $s_\ell$  is the number of units in the  $\ell^{th}$  layer **not** including the bias unit.

We want to minimise the cost by varying the components of *Theta*. To do this, we can use the **back propagation** algorithm. This is a generalisation of gradient descent which requires the partial derivatives  $\frac{\partial J(\Theta)}{\partial \theta_{ij}^\ell}$ .

## **Autonomous driving as an example:**

<https://www.coursera.org/learn/machine-learning/lecture/zYS8T/autonomous-driving>

## **Autonomous driving as an example:**

<https://www.coursera.org/learn/machine-learning/lecture/zYS8T/autonomous-driving>

Take a look at Python tutorial 4